

# PostgreSQL at Scale: Database Schema Changes with Minimal/No-Downtime

## Table of Contents

BRIEF OVERVIEW: .....	2
WHY SCHEMAS CHANGES ARE REQUIRED: .....	2
SCHEMA CHANGES SCENARIOS: .....	2
RISKS INVOLVED IN SCHEMA CHANGES:.....	3
<i>Locking:</i> .....	3
<i>Data Integrity issues:</i> .....	5
<i>Inconsistencies in Data:</i> .....	5
<i>Performance Degradation:</i> .....	5
<i>Data Entry conflicts:</i> .....	6
MITIGATIONS .....	6
<i>Use lineage tracking</i> .....	6
<i>Run end-to-end tests and Monitoring Performance</i> .....	7
<i>Ensure Data Consistency</i> .....	8
CONSIDERATIONS DURING SCHEMA CHANGES: .....	8
<i>Testing and Performance Considerations:</i> .....	9
SIMULATING SCHEMA CHANGES.....	10
<i>Example-1</i> .....	10
TOOLS AND BEST PRACTICES PROCESS: .....	11
<i>Liquibase</i> .....	11
<i>Flyway</i> .....	11
<i>Blue green Deployment.</i> .....	12
<i>pgroll</i> .....	13
<i>pg-osc</i> .....	26
MONITORING .....	30
<i>Locks:</i> .....	30
<i>CPU/Memory:</i> .....	32
<i>Disk Space:</i> .....	32
LIVE DEMO .....	33
<i>Customer Case Studies</i> .....	33
BEST PRACTICES .....	37
CONCLUSION: .....	38
DISCLAIMER: .....	38
REFERENCES: .....	39

## **Brief Overview:**

Database schema migrations can be a double-edged sword. They are essential for keeping our systems up to date and in sync with evolving application requirements, but often come bundled with a set of challenges that can give even the most seasoned developers and database administrators a nightmare. There's no one tool or a process to be followed for a seamless, minimal/no-downtime and risk-free release. This session will walk-through different schema release scenarios cut across the different phases of Software Development Life Cycle (SDLC) with relevant tools and best practices for each of them with hands-on demo. The pillars for Schema release would be -

- Ease of use
- Seamless
- Rollback
- Minimal/No-Downtime
- Risk Free

## **Why Schemas changes are required:**

- Support new features or functionality in your application
- Defect fixes
- Improve system performance by optimizing data structures
- Enhance data security through better constraints
- Adapt to evolving business requirements

## **Schema Changes Scenarios:**

### **The following operations can cause downtime or errors:**

1. Adding a column with a non-null default value to an existing table with minimal downtime
2. Changing the data type of a column
3. Renaming a table
4. Removing a column
5. Add a column with unique constraint
6. Running Vacuum Full with minimal downtime using pg\_repack
7. Adding/Deleting an index

## Risks involved in Schema Changes:

### Locking:

#### **ACCESS EXCLUSIVE**

- Blocks ALL operations
- Used by: DROP TABLE, TRUNCATE, VACUUM FULL

#### **EXCLUSIVE**

- Blocks most writes and schema changes
- Used by: REFRESH MATERIALIZED VIEW

#### **SHARE ROW EXCLUSIVE**

- Blocks concurrent data changes
- Used by: CREATE TRIGGER

#### **SHARE**

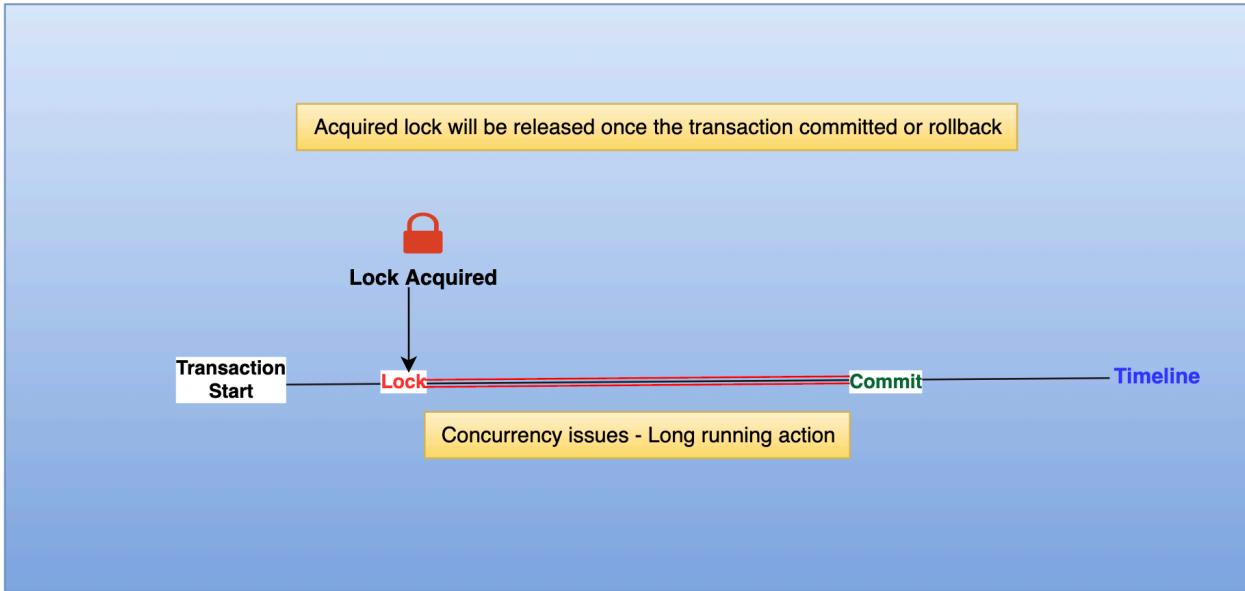
- Blocks structure changes
- Used by: CREATE INDEX (non-concurrent)

#### **ROW EXCLUSIVE**

- Allows concurrent reads
- Used by: UPDATE, DELETE, INSERT

For OLTP workloads (such as web and mobile applications), it is important to understand object-level and row-level locks in PostgreSQL. There are several good materials that are recommended for reading:

- The official documentation is a must-read, as usual: "[13.3. Explicit Locking](#)" (do not be confused by the title – this article discusses not only explicit locks that we can add using queries like `LOCK TABLE ...` or `SELECT ... FOR UPDATE` or functions like `pg_advisory_lock(...)` but also the locks introduced by regular SQL commands such as `ALTER` or `UPDATE`)
- "[PostgreSQL rocks, except when it blocks: Understanding locks](#)" (2018) by [Marco Slot](#) – a great extension to the docs, explaining some aspects and conveniently "translating" the table provided in the docs from "lock language" to "SQL command language"



The problem of holding exclusive locks for long can be very daunting. These can be locked rows (implicitly via `UPDATE` or `DELETE` or explicitly via `SELECT .. FOR UPDATE`) or database objects (example: successful `ALTER TABLE` inside a transaction block locks the table and holds the lock until the end of the transaction). If you need to learn more about locks in Postgres, read the article "[PostgreSQL rocks, except when it blocks: Understanding locks](#)" by Marco Slot. An abstract example of the general issue with locking:

```
begin;
alter table t1 add column c123 int8;
-- do something inside or outside of the database (or do nothing)
commit;
```

The reason for sitting inside a transaction after lock acquisition may vary. However, sometimes it is nothing – a simple waiting with an open transaction and acquired lock. This is the most annoying reason that can quickly lead to various performance or even partial downtime: an exclusive lock to a table blocks even `SELECT`s to this table.

#### **caution**

Remember: any lock acquired in a transaction is held until the very end of this transaction. It is released only when the transaction finishes, with either `COMMIT` or `ROLLBACK`.

**Every time we acquire an exclusive lock, we should think about finishing the transaction as soon as possible.**

### **Data Integrity issues:**

Systems often rely on specific schema structures to maintain data integrity. When you modify these structures, you risk creating situations where incoming data no longer meets your database's requirements. This can lead to failed insertions, constraint violations, and data quality issues.

For example, if you add a new required column without providing a default value, upstream systems might fail to insert new records. Similarly, changing data type constraints might cause previously valid data to be rejected. These issues can be particularly problematic in systems with real-time data ingestion requirements.

Schema changes can introduce subtle data inconsistencies, particularly when dealing with related tables or denormalized data structures.

### **Inconsistencies in Data:**

For example, if you modify a column in one table but forget to update corresponding columns in related tables, you might create mismatches that negatively affect data integrity. It's bad news when inconsistencies pop up in complex systems or those relying on materialized views. The web of relationships quickly becomes overwhelming.

The ripple effects of these inconsistencies often surface in join operations, where mismatched data types or column names can lead to incorrect results or missing data. These issues might not be immediately apparent, only popping up for specific edge cases or when reconciling data across different systems.

### **Performance Degradation:**

Tweaking indexes or data types can have serious consequences on performance. Even a seemingly simple change - like adding a new column with a default value to a large table - can cause performance degradation during the modification process. In addition, there could be cascading impact that spreads far and wide. Queries that previously performed well might suddenly slow down as the query optimizer chooses different plans based on the modified schema.

#### a) Table-level Locking:

- ACCESS EXCLUSIVE locks blocking reads/writes
- Connection pileup due to lock waiting

- Transaction queue buildup

#### b) I/O Impact:

- Increased disk writes during table rewrites
- Buffer cache pollution
- WAL generation spike

#### **Data Entry conflicts:**

When schema changes aren't properly coordinated with up/down stream data providers, you risk creating misalignments between what these systems are sending and what your database expects to receive. In systems where data pours in from multiple sources, it's trial by fire to manage the flow.

Common data entry issues include:

- Forms and applications submitting data in outdated formats
- ETL processes using incorrect field mappings
- API integrations failing to adapt to new requirements
- Batch processes sending incompatible data structures

Common data entry issues include:

- Forms and applications submitting data in outdated formats
- ETL processes using incorrect field mappings
- API integrations failing to adapt to new requirements
- Batch processes sending incompatible data structures

## **Mitigations**

### **Use lineage tracking**

One of the most effective ways to manage the impacts of schema changes is through data lineage tracking. Data lineage tools can automatically track these relationships, making it easier to assess the potential impact of any proposed changes.

By maintaining a clear map of table-to-table dependencies and column-level relationships, you can better understand how changes will propagate through your system.

Effective lineage tracking should go beyond documenting direct dependencies. It should capture transformation logic, business rules, and data quality requirements. This holistic view enables you to predict how schema changes might affect not just the immediate consumers of your data but also downstream processes and business operations.

## Run end-to-end tests and Monitoring Performance

Testing schema changes requires a systematic approach that goes beyond simple functional verification. You need to establish a comprehensive testing strategy that includes regression testing, performance testing, and validation of business logic. This testing should occur in an environment that closely mirrors production, with realistic data volumes and usage patterns.

You can't afford to get bogged down in manual testing. Automated tools are the key to streamlining your process, weeding out hiccups, and pushing through to launch.

- Performance testing under various load conditions
- Integration testing with downstream systems
- Validation of business rules and calculations
- Testing of edge cases and error conditions

These should include:

- Data quality validation to verify that transformations maintain data integrity
- Performance testing under various load conditions
- Integration testing with downstream systems
- Validation of business rules and calculations
- Testing of edge cases and error conditions

Performance monitoring should be implemented before, during, and after schema changes. Numbers to watch include query execution times, resource usage, and latency. Tracking your baseline performance metrics before making adjustments will help you instantly spot any areas that need attention.

## Performance Optimization Techniques

### a) Memory Configuration:

```
# Adjust work_mem for schema changes
work_mem = 64MB
maintenance_work_mem = 256MB
```

### b) Autovacuum Settings:

```
autovacuum_vacuum_scale_factor = 0.1  
autovacuum_analyze_scale_factor = 0.05
```

## Implementation Strategies

### a) Pre-migration Optimization:

```
-- Analyze table before changes  
VACUUM ANALYZE table_name;  
  
-- Remove dead tuples  
VACUUM FULL table_name;
```

## Ensure Data Consistency

Maintaining data consistency during schema changes requires a multi-faceted approach. Before you begin, assess the connections between affected tables and columns, considering both the rules set in stone and the unwritten rules of the business. You need to be surgical when making changes to your data model to avoid downstream failures.

For example, if you modify a customer identifier format in one table, you need to make sure that this change is reflected in all related tables and views. This might involve creating temporary staging tables, implementing parallel processing strategies, or using transitional periods where both old and new formats are supported. It's time-consuming, but being thorough about proactively maintaining data quality and consistency can prevent even bigger problems down the line.

### Few other Mitigations:

- Deploy changes during low-traffic window and have a rollback plan
- Use staging environments
- Implement versioning
- Communicate with stakeholders

## Considerations during Schema Changes:

**Plan for Rollbacks:** Always have a rollback plan in place. If a migration does not go

as planned, being able to revert to the previous state quickly can save time and prevent user disruption

- Maintain old structure until verification
- Keep backup of original indexes
- Monitor application performance metrics

**Monitor Performance:** Schema changes can be resource -intensive, Monitoring is crucial to prevent performance issues. After deploying schema changes, monitor the performance of the application closely. Look for any signs of locking or slow queries that may indicate issues stemming from the migration.

#### **Ensure Backward Compatibility:**

When making changes, ensure that both old and new code can run simultaneously. This is crucial in environments where both versions of the application are operational, as breaking changes can lead to outages.

#### Other Important Considerations:

1. Always backup before schema changes
2. Test migrations in staging environment
3. Consider impact on foreign keys and constraints
4. Plan for adequate maintenance window

By following these practices and understanding the potential inconsistencies, you can minimize risks during online schema changes while maintaining data integrity and application availability.

Remember that PostgreSQL 15+ offers improved concurrent DDL operations, making some of these problems easier to handle, but careful planning is still essential for production environments.

#### **Testing and Performance Considerations:**

**Query Analysis:** Utilize `pg_stat_statements` to identify hot or slow queries, which can help in optimizing performance.

**Load Testing:** Conduct load testing preferably on a staging environment. Tools like [k6](#) can simulate traffic from multiple users, allowing you to assess how your database performs under stress.

**Traffic Surge Preparation:** Anticipate surges in traffic, especially during significant launches.

## Simulating Schema Changes

Below is a schema change use-case illustrating the impacts on the system from performance and resource contention standpoint

## Example-1

Add column to the transactional table which will be almost busy during the business hours.

## Impacts:

## **Table Locking**

Depending on the database system, adding a column might lock the entire table. This prevents other transactions from accessing the table and can cause application timeouts and business disruptions

## ***Performance Impact***

- Heavy CPU usage during the operation
  - Increased I/O operations
  - Slower response times for other queries
  - Potential system-wide performance degradation

```

top - 10:49:55 up 29 days, 4:44, 10 users, load average: 20.33, 16.56, 10.05
Tasks: 193 total, 24 running, 169 sleeping, 0 stopped, 0 zombie
%Cpu(s): 14.7 us, 11.8 sy, 0.0 ni, 4.9 id, 66.6 wa, 0.0 hi, 1.2 si, 0.8 st
MiB Mem : 3904.2 total, 516.5 free, 348.7 used, 3039.0 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used. 3154.5 avail Mem

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1648745	postgres	20	0	381124	5480	4656	S	10.0	0.1	3:13.22	pgbench
155	root	0	-20	0	0	0	I	4.3	0.0	0:35.52	kworker/1:1H-kblockd
46	root	0	-20	0	0	0	I	1.7	0.0	0:27.91	kworker/0:1H-kblockd
1647125	postgres	20	0	427480	17284	14412	S	1.7	0.4	0:28.67	postgres
1648754	postgres	20	0	427252	151832	149024	R	1.3	3.8	0:19.67	postgres
1648755	postgres	20	0	427252	151876	149036	R	1.3	3.8	0:19.27	postgres
1648756	postgres	20	0	427252	151548	148744	R	1.3	3.8	0:19.29	postgres

## Space Issues

Immediate storage requirement if adding with default values, table size increase might cause storage problems

## Tools and Best Practices Process:

### Liquibase

[Liquibase](#) is an open-source database schema change management and version control tool that helps track, version, and deploy database changes. It's particularly useful in DevOps environments and for teams working on applications that require database modifications.

### Key Features of Liquibase:

- Version Control
- Tracks database changes through change sets
- Maintains a history of all database modifications
- Enables rollback capabilities

### Format Supported

- XML
- YAML
- JSON
- SQL
- Format-specific change log files

### Flyway

[Flyway](#) is a database migration tool that helps manage database schema changes

across environments. It works by applying versioned migrations in order, tracking what has been applied in a special table called "flyway\_schema\_history".

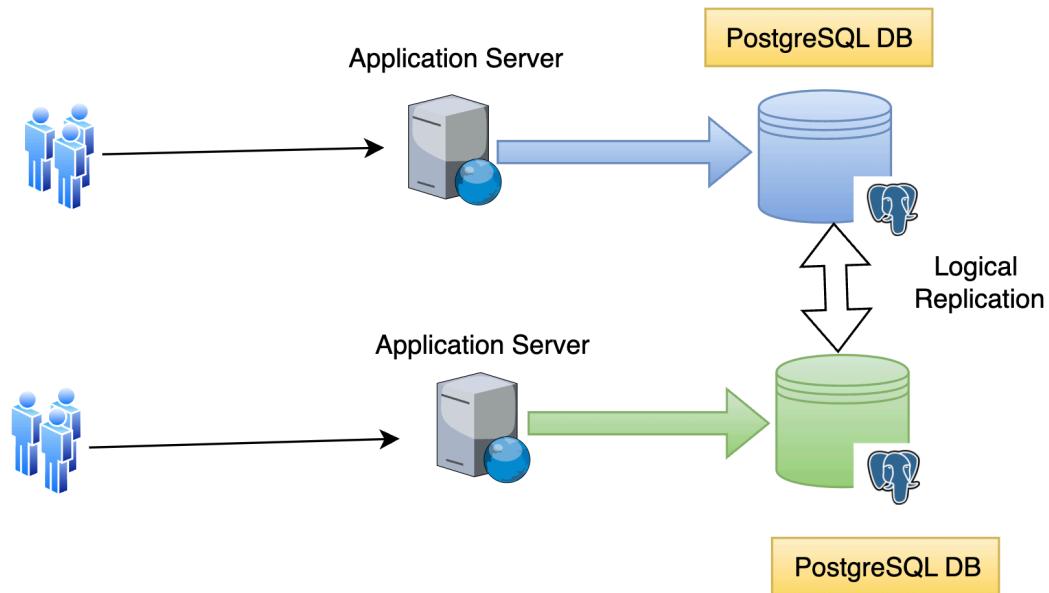
### Key Features:

- Simple version-based migrations
- SQL and Java-based migrations
- Command-line and API support
- Easy integration with build tools
- Schema history tracking

### Blue green Deployment

Blue-green deployment is an application release strategy that uses two identical production environments — referred to as “blue” (live) and “green” (standby) — to achieve reliable testing, zero-downtime releases, and instant rollbacks. At any given time, only one of these environments is live.

For example, the blue environment handles all live production traffic while the green environment remains idle. When a new application version is ready, it is deployed to the green environment for testing. Once the release passes testing, traffic is seamlessly rerouted from blue to green.



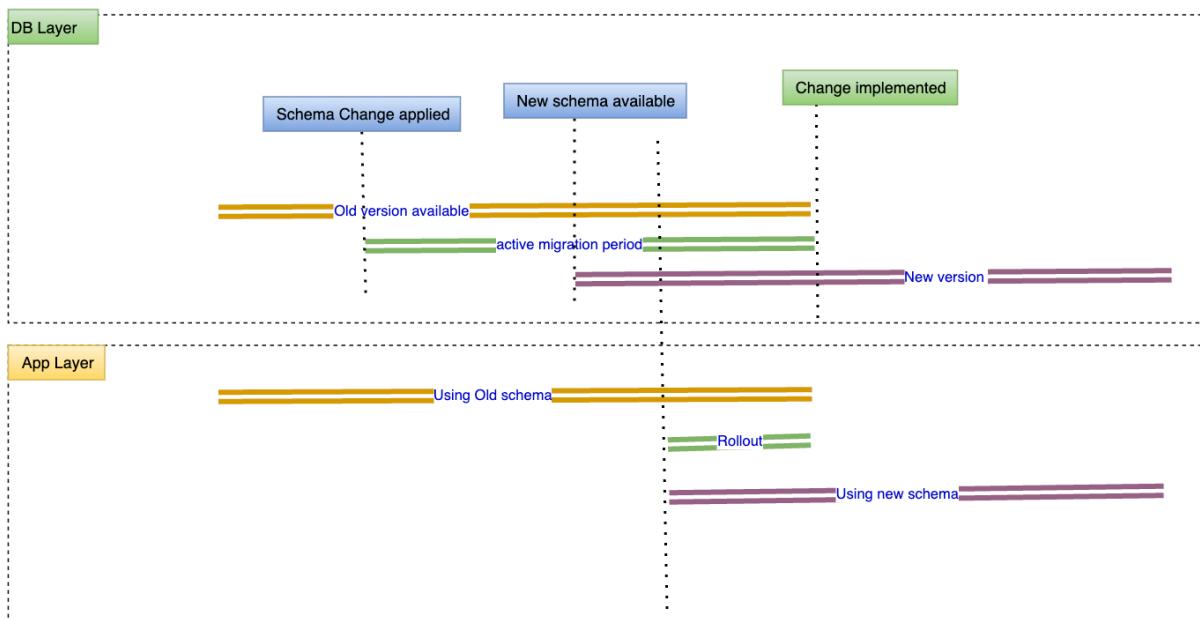
At this point, the green environment becomes the new live production environment, and the blue environment goes idle. This deployment strategy ensures that any issues

that arise during or after the transition can be easily mitigated by reverting traffic back to the previous (blue) version.

## pgroll

[PgRoll](#) is a specialized database schema migration tool designed specifically for PostgreSQL, focusing on zero-downtime migrations. Created by Xata, it addresses the critical need for performing complex database changes without service interruption. Unlike traditional migration tools, pgroll implements sophisticated strategies to maintain database availability while ensuring data consistency throughout the migration process.

The tool represents a significant advancement in database DevOps, particularly for organizations running mission-critical applications where downtime is not acceptable. It combines the reliability of traditional migration tools with modern zero-downtime requirements, making it invaluable for continuous deployment environments.



## How to setup pgroll

### Pre-Requisites:

To install the pgroll go language should be installed in the server whose version should be 1.23 and above.

### Steps to install go and pgroll

1. Download appropriate sources/binaries need compatible with your OS from below link:
  - a. <https://go.dev/doc/install> and follow steps present in the link to install go.
2. Once the go is installed install pgroll

```
mkdir /opt/pgroll
cd /opt/pgroll
go install github.com/xataio/pgroll@latest
```

3. Check the connection with PostgreSQL and install initial schema required for pgroll

```
pgroll init --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable
```

```
[postgres@ip-172-31-17-243 ~]$ pgroll init --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable
SUCCESS Initialization complete
[postgres@ip-172-31-17-243 ~]$
[postgres@ip-172-31-17-243 ~]$
[postgres@ip-172-31-17-243 ~]$
[postgres@ip-172-31-17-243 ~]$
[postgres@ip-172-31-17-243 ~]$
[postgres@ip-172-31-17-243 ~]$ psql
psql (15.9)
Type "help" for help.

postgres=# \c pgconf
You are now connected to database "pgconf" as user "postgres".
pgconf=# \dn
          List of schemas
   Name  |    Owner
-----+-----
 pgroll | postgres
 public | pg_database_owner
(2 rows)
```

**Note:** In case if pgroll is not in PATH then set the PATH in .bash\_profile for Postgres user.

pgroll will create its own schema and in it there will be table by name “migrations” which will hold the information of all tables which have entry of all changes

```
expanded_triggers is off
pgconf# select * from migrations limit 1;
-[ RECORD 1 ]-----+
schema      | public_02_change_column_type
name        | sql_39b9eb04edb340
migration   | {"name": "sql_39b9eb04edb340", "operations": [{"sql": {"up": "drop schema public_02_change_column_type cascade;"} }]}
created_at  | 2025-01-22 07:10:19.748752
updated_at  | 2025-01-22 07:10:19.748752
parent      |
done        | t
resulting_schema | {"name": "public_02_change_column_type", "tables": {}}
migration_type | inferred
```

## ***Example Schema Change Scenarios using pgroll***

### **1. Adding a column with a not-null default value to an existing table with minimal downtime**

#### **Add column**

An add column operation creates a new column on an existing table.

#### **Structure**

```
{  
  "add_column": {  
    "table": "name of table to which the column should be  
added",  
    "up": "SQL expression",  
    "column": {  
      "name": "name of column",  
      "type": "postgres type",  
      "comment": "postgres comment for the column",  
      "nullable": true|false,  
      "unique": true|false,  
      "pk": true|false,  
      "default": "default value for the column",  
      "check": {  
        "name": "name of check constraint",  
        "constraint": "constraint expression"  
      },  
      "references": {  
        "name": "name of foreign key constraint",  
        "table": "name of referenced table",  
        "column": "name of referenced column",  
        "on_delete": "ON DELETE behaviour, can be CASCADE, SET  
NULL, RESTRICT, or NO ACTION. Default is NO ACTION",  
      }  
    }  
  }  
}
```

The up SQL is used when a row is added to the old version of the table (ie, without the new column). In the new version of the table, the row's value for the new column is determined by the up SQL expression. The up SQL can be omitted if the new column is nullable or has a DEFAULT defined.

Default values are subject to the usual rules for quoting SQL expressions. In particular, string literals should be surrounded with single quotes.

**NOTE:** As a special case, the up field can be omitted when adding smallserial, serial and bigserial columns.

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    name VARCHAR(100)
);
```

```
ALTER TABLE users
ADD COLUMN email VARCHAR(255);
```

*Example:*

```
vi 01_add_column_to_products.json
```

```
{
  "name": "01_add_column_to_products",
  "operations": [
    {
      "add_column": {
        "table": "users",
        "column": {
          "name": "email",
          "type": "varchar(255)",
          "nullable": true
        }
      }
    },
    {
      "add_column": {
        "table": "users",
        "column": {
          "name": "cost",
          "type": "int",
          "nullable": true
        }
      }
    }
  ]
}
```

## *Execution of schema change*

```
pgroll start --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable
01_add_column_to_products.json
pgroll complete --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable
01_add_column_to_products.json
```

```
[postgres@ip-172-31-17-243 ~]$ pgroll start --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 01_add_column_to_products.json
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf
pgsql (15.9)
Type "help" for help.

pgconf=# \dn
           List of schemas
 Schema | Owner
-----+
 pgcoll | postgres
 public  | pg_database_owner
 public_01_add_column_to_products | postgres
(3 rows)

pgconf=# \dt+ users
          Table "public.users"
 Column | Type | Collation | Nullable | Default | Storage | Compression | Stats target | Description
-----+-----+-----+-----+-----+-----+-----+-----+-----+
 id    | integer | not null |          | nextval('users_id_seq'::regclass) | plain   | extended   |             |
 name  | character varying(100) |          |          |          | plain   | extended   |             |
 _pgsql_new_email | character varying(255) |          |          |          | plain   | extended   |             |
 _pgsql_new_cost  | integer |          |          |          | plain   |             |             |
Indexes:
 "users_pkey" PRIMARY KEY, btree (id)
Access method: heap
pgconf=#
```

```
[postgres@ip-172-31-17-243 ~]$ pgroll complete --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 01_add_column_to_products.json
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf
pgsql (15.9)
Type "help" for help.

pgconf=# \q
[postgres@ip-172-31-17-243 ~]$ pgconf
pgsql (15.9)
Type "help" for help.

pgconf=# \dt+ users
          Table "public.users"
 Column | Type | Collation | Nullable | Default | Storage | Compression | Stats target | Description
-----+-----+-----+-----+-----+-----+-----+-----+-----+
 id    | integer | not null |          | nextval('users_id_seq'::regclass) | plain   | extended   |             |
 name  | character varying(100) |          |          |          | plain   | extended   |             |
 email | character varying(255) |          |          |          | plain   | extended   |             |
 cost  | integer |          |          |          | plain   |             |             |
Indexes:
 "users_pkey" PRIMARY KEY, btree (id)
Access method: heap
```

## 2. Changing the data type of a column

### *Change Datatype*

A **change type** operation changes the type of a column.

*Structure*

```
{ "alter_column": { "table": "table name", "column": "column name", "type": "new type of column", }
```

```

        "up": "SQL expression",
        "down": "SQL expression"
    }
}

```

Use the up SQL expression to do data conversion from the old column type to the new type. In the old schema version, the column will have its old data type; in the new version the column will have its new type.

Use the down SQL expression to do data conversion in the other direction; from the new data type back to the old.

*Example:*

```

ALTER TABLE users
alter COLUMN type bigint;

vi 02_change_column_type.json

```

```

{
  "name": "02_change_column_type",
  "operations": [
    {
      "alter_column": {
        "table": "users",
        "column": "cost",
        "type": "float",
        "up": "float",
        "down": "integer"
      }
    }
  ]
}

```

*Execution of schema change*

```

pgroll start --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disab
le 02_change_column_type.json
pgroll complete --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disab
le 02_change_column_type.json

```

```
[postgres@ip-172-31-17-243 ~]$ pgroll start --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 02_change_column_type.json
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf
psql (15.9)
Type "help" for help.

pgconf=# \d+ users
          Column          |      Type       | Collation | Nullable | Default | Table "public.users" | Storage | Compression | Stats target | Description
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
 id   | integer      |           | not null | nextval('users_id_seq'::regclass) |           | plain    | extended    |             | 
 name | character varying(100) |           |           |           |           | plain    | extended    |             | 
 email | character varying(255) |           |           |           |           | plain    | extended    |             | 
 cost  | integer      |           |           |           |           | plain    | plain       |             | 
_pgroll_new_cost | double precision |           |           |           |           | plain    | plain       |             | 
Indexes:
"users_pkey" PRIMARY KEY, btree (id)
Triggers:
_pgroll_trigger_users_pgroll_new_cost BEFORE INSERT OR UPDATE ON users FOR EACH ROW EXECUTE FUNCTION _pgroll_trigger_users_pgroll_new_cost()
_pgroll_trigger_users_cost BEFORE INSERT OR UPDATE ON users FOR EACH ROW EXECUTE FUNCTION _pgroll_trigger_users_cost()
Access method: heap

pgconf=# \q
[postgres@ip-172-31-17-243 ~]$ pgroll complete --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 02_change_column_type.json
SUCCESS Migration successful!
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf
psql (15.9)
Type "help" for help.

pgconf=# \d+ users
          Column          |      Type       | Collation | Nullable | Default | Table "public.users" | Storage | Compression | Stats target | Description
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
 id   | integer      |           | not null | nextval('users_id_seq'::regclass) |           | plain    | extended    |             | 
 name | character varying(100) |           |           |           |           | plain    | extended    |             | 
 email | character varying(255) |           |           |           |           | plain    | extended    |             | 
 cost  | double precision |           |           |           |           | plain    | plain       |             | 
Indexes:
"users_pkey" PRIMARY KEY, btree (id)
Access method: heap
```

### 3. Renaming a table

#### *Rename table*

A rename table operation renames a table.

#### *Structure*

```
{
  "rename_table": {
    "from": "old column name",
    "to": "new column name"
  }
}
```

The table is accessible by its old name in the old version of the schema, and by its new name in the new version of the schema.

The table itself is renamed on migration completion.

#### *Example:*

```
pgconf=# create table customers(id int primary key, cust_name
text,cust_add text);
CREATE TABLE
pgconf=# \d customers
Table "public.customers"
Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----+
```

```
id | integer | | not null |
cust_name | text | | |
cust_add | text | | |
Indexes:
"customers_pkey" PRIMARY KEY, btree (id)
```

```
pgconf=# \q
[postgres@ip-172-31-17-243 ~]$
```

```
vi 03_rename_table.json
{
  "name": "03_rename_table",
  "operations": [
    {
      "rename_table": {
        "from": "customers",
        "to": "clients"
      }
    }
  ]
}
```

### *Execution of schema change*

```
pgroll start --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable
03_rename_table.json
pgroll complete --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable
03_rename_table.json
```

```
[postgres@ip-172-31-17-243 ~]$ pgroll start --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 03_rename_table.json
[SUCCESS] New version of the schema available under the postgres "public_03_rename_table" schema
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf
psql (15.9)
Type 'help' for help.

pgconf=# \dt
      List of relations
 Schema |   Name    | Type  | Owner
-----+-----+-----+-----+
 public | customers | table | postgres
 public | users     | table | postgres
(2 rows)

pgconf=# \d customers
          Table "public.customers"
 Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----+
 id    | integer |           | not null |
 cust_name | text |           |           |
 cust_addr | text |           |           |
Indexes:
 "customers_pkey" PRIMARY KEY, btree (id)
Access method: heap

pgconf=# \d+ customers
          Table "public.customers"
 Column | Type | Collation | Nullable | Default | Storage | Compression | Stats target | Description
-----+-----+-----+-----+-----+-----+-----+-----+-----+
 id    | integer |           | not null |           | plain   | extended   |           |
 cust_name | text |           |           |           |           |           |           |
 cust_addr | text |           |           |           |           |           |           |
Indexes:
 "customers_pkey" PRIMARY KEY, btree (id)
Access method: heap

pgconf=# \q
[postgres@ip-172-31-17-243 ~]$ pgroll complete --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 03_rename_table.json
[SUCCESS] Migration successful!
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf -c "\dt"
      List of relations
 Schema |   Name    | Type  | Owner
-----+-----+-----+-----+
 public | clients  | table | postgres
 public | users    | table | postgres
(2 rows)
```

## 4. Removing a column

### *Drop column*

A drop column operation drops a column from an existing table.

#### *Structure*

```
{
  "drop_column": {
    "table": "name of table",
    "column": "name of column to drop",
    "down": "SQL expression"
  }
}
```

The `down` field above is required in order to back fill the previous version of the schema during an active migration.

#### *Example:*

```
vi 04_drop_column.json
{
  "name": "04_drop_column",
  "operations": [
```

```
{
  "drop_column": {
    "table": "users",
    "column": "cost",
    "down": "0"
  }
}
]
```

### *Execution of schema change*

```
pgroll start --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 04_drop_column.json
pgroll complete --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 04_drop_column.json
```

```
[postgres@ip-172-31-17-243 ~]$ pgroll start --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 04_drop_column.json
[SUCCESS] New version of the schema available under the postgres "public_04_drop_column" schema
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf -c "\d users"
Table "public.users"
 Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----+
 id   | integer |           | not null | nextval('users_id_seq'::regclass)
 name | character varying(100) |
 email | character varying(255) |
 cost  | double precision |
Indexes:
 "users_pkey" PRIMARY KEY, btree (id)
Triggers:
 _pgroll_trigger_users_cost BEFORE INSERT OR UPDATE ON users FOR EACH ROW EXECUTE FUNCTION _pgroll_trigger_users_cost()
[postgres@ip-172-31-17-243 ~]$ pgroll complete --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 04_drop_column.json
[SUCCESS] Migration successful!
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf -c "\d users"
Table "public.users"
 Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----+
 id   | integer |           | not null | nextval('users_id_seq'::regclass)
 name | character varying(100) |
 email | character varying(255) |
Indexes:
 "users_pkey" PRIMARY KEY, btree (id)
```

### Add a column with unique constraint

#### *Add unique constraint*

Add unique operations add a 'UNIQUE' constraint to a column.

#### *Structure*

```
{
  "alter_column": {
    "table": "table name",
    "column": "column name",
    "unique": {
      "name": "name of unique constraint"
    }
}
```

```

        },
        "up": "SQL expression",
        "down": "SQL expression"
    }
}

```

Use the `up` SQL expression to migrate values from the old non-unique column in the old schema to the `UNIQUE` column in the new schema.

*Example:*

```
vi 06_create_tickets_table.json
```

```
{
  "name": "06_create_tickets_table",
  "operations": [
    {
      "create_table": {
        "name": "tickets",
        "columns": [
          {
            "name": "ticket_id",
            "type": "serial",
            "pk": true
          },
          {
            "name": "sellers_name",
            "type": "varchar(255)"
          },
          {
            "name": "sellers_zip",
            "type": "integer"
          },
          {
            "name": "ticket_type",
            "type": "varchar(255)",
            "default": "'paper'"
          }
        ]
      }
    }
  ]
}
```

## *Execution of schema change*

```
pgroll start --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_create_tickets_table
pgroll complete --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_create_tickets_table
```

```
[postgres@ip-172-31-17-243 ~]$ pgroll start --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_create_tickets_table.json
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf
psql (15.9)
Type "help" for help.

pgconf=# \dt
      List of relations
 Schema |   Name    | Type | Owner
-----+-----+-----+-----+
 public | _pgroll_new_tickets | table | postgres
 public | clients        | table | postgres
 public | users          | table | postgres
(3 rows)

pgconf=# \dn
      List of schemas
 Name | Owner
-----+-----
 pgroll | postgres
 public | pg_database_owner
 public_04_drop_column | postgres
 public_06_create_tickets_table | postgres
(4 rows)

pgconf=# exit
[postgres@ip-172-31-17-243 ~]$ pgroll complete --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_create_tickets_table.json
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf -c "\dt"
      List of relations
 Schema |   Name    | Type | Owner
-----+-----+-----+-----+
 public | clients | table | postgres
 public | tickets | table | postgres
 public | users  | table | postgres
(3 rows)
```

```
vi 06_01_add_table_unique_constraint.json
{
  "name": "06_01_add_table_unique_constraint",
  "operations": [
    {
      "create_constraint": {
        "type": "unique",
        "table": "tickets",
        "name": "unique_zip_name",
        "columns": [
          "sellers_name",
          "sellers_zip"
        ],
        "up": {
          "sellers_name": "sellers_name",
          "sellers_zip": "sellers_zip"
        },
        "down": {
          "sellers_name": "sellers_name",
          "sellers_zip": "sellers_zip"
        }
      }
    }
  ]
}
```

```

        "sellers_name": "sellers_name",
        "sellers_zip": "sellers_zip"
    }
}
]
}

pgroll start --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_01_add_table_unique_constraint.json
pgroll complete --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_01_add_table_unique_constraint.json

```

```

[postgres@ip-172-31-17-243 ~]$ psql -d pgconf
psql (15.3)
Type 'help' for help.

pgconf# \d tickets
                                         Table "public.tickets"
  Column   | Type      | Collation | Nullable | Default
-----+-----+-----+-----+-----+
ticket_id | integer   |           | not null | nextval('_pgroll_new_tickets_ticket_id_seq)::regclass'
sellers_name | character varying(255) |           | not null |
sellers_zip | integer   |           | not null |
ticket_type | character varying(255) |           | not null | 'paper'::character varying
Indexes:
  "_pgroll_new_tickets_pkey" PRIMARY KEY, btree (ticket_id)

pgconf# \q
[postgres@ip-172-31-17-243 ~]$ pgroll start --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_01_add_table_unique_constraint.json
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf -c "\d+ tickets"
                                         Table "public.tickets"
  Column   | Type      | Collation | Nullable | Default          | Storage | Compression | Stats target | Description
-----+-----+-----+-----+-----+-----+-----+-----+-----+
ticket_id | integer   |           | not null | nextval('_pgroll_new_tickets_ticket_id_seq)::regclass | plain   |           |           |
sellers_name | character varying(255) |           | not null |                               | extended |           |           |
sellers_zip | integer   |           | not null |                               | plain   |           |           |
ticket_type | character varying(255) |           | not null | 'paper'::character varying | extended |           |           |
_pgroll_new_sellers_name | character varying(255) |           |           |                               | extended |           |           |
_pgroll_new_sellers_zip | integer   |           |           |                               | plain   |           |           |
Indexes:
  "_pgroll_new_tickets_pkey" PRIMARY KEY, btree (ticket_id)
  "_unique_zip_name" UNIQUE, btree (_pgroll_new_sellers_name, _pgroll_new_sellers_zip)
Check constraints:
  "_pgroll_dup_pgroll_check_not_null_sellers_name" CHECK (_pgroll_new_sellers_name IS NOT NULL) NOT VALID
  "_pgroll_dup_pgroll_check_not_null_sellers_zip" CHECK (_pgroll_new_sellers_zip IS NOT NULL) NOT VALID
Triggers:
  _pgroll_trigger_tickets_pgroll_new_sellers_name BEFORE INSERT OR UPDATE ON tickets FOR EACH ROW EXECUTE FUNCTION _pgroll_trigger_tickets_pgroll_new_sellers_name()
  _pgroll_trigger_tickets_pgroll_new_sellers_zip BEFORE INSERT OR UPDATE ON tickets FOR EACH ROW EXECUTE FUNCTION _pgroll_trigger_tickets_pgroll_new_sellers_zip()
  _pgroll_trigger_tickets_sellers_name BEFORE INSERT OR UPDATE ON tickets FOR EACH ROW EXECUTE FUNCTION _pgroll_trigger_tickets_sellers_name()
  _pgroll_trigger_tickets_sellers_zip BEFORE INSERT OR UPDATE ON tickets FOR EACH ROW EXECUTE FUNCTION _pgroll_trigger_tickets_sellers_zip()
Access method: heap

[postgres@ip-172-31-17-243 ~]$ pgroll complete --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_01_add_table_unique_constraint.json
[SUCCESS] Migration successful!
[postgres@ip-172-31-17-243 ~]$ psql -d pgconf -c "\d+ tickets"
                                         Table "public.tickets"
  Column   | Type      | Collation | Nullable | Default          | Storage | Compression | Stats target | Description
-----+-----+-----+-----+-----+-----+-----+-----+-----+
ticket_id | integer   |           | not null | nextval('_pgroll_new_tickets_ticket_id_seq)::regclass | plain   |           |           |
ticket_type | character varying(255) |           | not null | 'paper'::character varying | extended |           |           |
sellers_name | character varying(255) |           | not null |                               | extended |           |           |
sellers_zip | integer   |           | not null |                               | plain   |           |           |
Indexes:
  "_pgroll_new_tickets_pkey" PRIMARY KEY, btree (ticket_id)
  "_unique_zip_name" UNIQUE CONSTRAINT, btree (sellers_name, sellers_zip)
Access method: heap

```

Other command line options are available to rollback, check the status

## **Rollback**

Roll back the currently active migration.

### ***Command***

```

$ pgroll rollback --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_01_add_table_unique_constraint.json

```

This rolls back the currently active migration (an active migration is one that has been started but not yet completed).

Rolling back a pgroll migration means removing the new schema version. The old schema version was still present throughout the migration period and does not require modification.

Migrations cannot be rolled back once completed. Attempting to roll back a migration that has already been completed is a no-op.

## Status

Show the current status of pgroll within a given schema.

### Command

```
$ pgroll status --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_01_add_table_unique_constraint.json
```

The status field can be one of the following values:

- "No migrations" - no migrations have been applied in this schema yet.
- "In progress" - a migration has been started, but not yet completed.
- "Complete" - the most recent migration was completed.

The Version field gives the name of the latest schema version.

If a migration is In progress the schemas for both the latest version indicated by the Version field and the previous version will exist in the database.

If a migration is Complete only the latest version of the schema will exist in the database.

The top-level --schema flag can be used to view the status of pgroll in a different schema:

```
$ pgroll status --postgres-url
postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_01_add_table_unique_constraint.json --schema public
```

```
postgres@ip-172-31-17-243 ~]$ pgroll status --postgres-url postgres://postgres:postgres@localhost:5432/pgconf?sslmode=disable 06_01_add_table_unique_constraint.json --schema public
{
  "schema": "public",
  "version": "06_01_add_table_unique_constraint",
  "status": "Complete"
}
```

For more details on [pgroll](#)

## pg-osc

[pg\\_osc](#), or PostgreSQL Online Schema Change – an *open-source tool* designed to revolutionize the execution of schema modifications without imposing substantial downtime or locking the entire database. This tool proves invaluable for making

crucial adjustments to the database structure, ranging from adding or removing columns to changing data types and creating or dropping indexes.

## Pre-Requisites

1. Clone the pg-osc repository and build extension:

```
yum install ruby* gem* -y  
gem update --system 3.6.3
```

```
git clone https://github.com/shayonj/pg-osc.git  
  
cd pg-osc  
gem install pg_online_schema_change
```

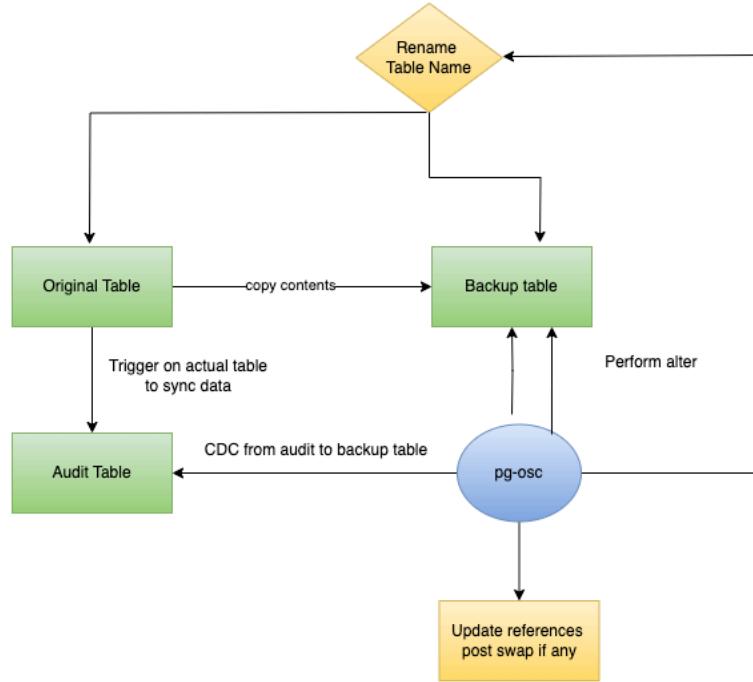
```
[root@ip-172-31-17-243 pg-osc]# gem install pg_online_schema_change  
Successfully installed google-protobuf-3.25.5-x86_64-linux  
Building native extensions. This could take a while...  
Successfully installed pg_query-4.2.3  
Fetching pg-1.5.9.gem  
Building native extensions. This could take a while...  
Successfully installed pg-1.5.9  
Fetching ostruct-0.6.1.gem  
Successfully installed ostruct-0.6.1  
Fetching bigdecimal-3.1.9.gem  
Building native extensions. This could take a while...  
Successfully installed bigdecimal-3.1.9  
Fetching oj-3.16.9.gem  
Building native extensions. This could take a while...  
Successfully installed oj-3.16.9  
Fetching ougai-2.0.0.gem  
Successfully installed ougai-2.0.0  
Fetching pg_online_schema_change-0.9.10.gem  
Successfully installed pg_online_schema_change-0.9.10  
Parsing documentation for google-protobuf-3.25.5-x86_64-linux  
Installing ri documentation for google-protobuf-3.25.5-x86_64-linux  
Parsing documentation for pg_query-4.2.3  
Installing ri documentation for pg_query-4.2.3  
Parsing documentation for pg-1.5.9  
Installing ri documentation for pg-1.5.9  
Parsing documentation for ostruct-0.6.1  
Installing ri documentation for ostruct-0.6.1  
Parsing documentation for bigdecimal-3.1.9  
Installing ri documentation for bigdecimal-3.1.9  
Parsing documentation for oj-3.16.9  
Installing ri documentation for oj-3.16.9  
Parsing documentation for ougai-2.0.0  
Installing ri documentation for ougai-2.0.0  
Parsing documentation for pg_online_schema_change-0.9.10  
Installing ri documentation for pg_online_schema_change-0.9.10  
Done installing documentation for pg_online_schema_change after 9 seconds  
8 gems installed  
[root@ip-172-31-17-243 pg-osc]#
```

```
-rwxr--r-- 1 root root 131 Jan 29 06:33 Maxfile  
[root@ip-172-31-17-243 pg-osc]# cd bin/  
[root@ip-172-31-17-243 bin]# ls -lr  
total 12  
-rwxr--r-- 1 root root 124 Jan 29 06:33 pg-online-schema-change  
-rwxr--r-- 1 root root 134 Jan 29 06:33 console  
[root@ip-172-31-17-243 bin]# ./pg-online-schema-change --help  
Commands:  
  pg-online-schema-change --version -v                                     # print the version  
  pg-online-schema-change help [COMMAND]                                # describe available commands or one specific command  
  pg-online-schema-change perform -a, --alter-statement=ALTER_STATEMENT -d, --dbname=DBNAME -h, --host=HOST -p, --port=N -s, --schema=SCHEMA -u, --username=USERNAME # safely apply schema changes with minimal locks
```

Set path in the .bash\_profile for PostgreSQL user

## How does it work

- **Primary table:** A table against which a potential schema change is to be run
- **Shadow table:** A copy of an existing primary table
- **Audit table:** A table to store any updates/inserts/delete on a primary table



## Use case for pg-osc

### *Multiple ALTER statements*

```
pg-online-schema-change perform --alter-statement 'ALTER TABLE
users ADD COLUMN "gender" BOOLEAN DEFAULT FALSE; ALTER TABLE
users ADD COLUMN "last_name" varchar(20);' --dbname "pgconf" --
host "localhost" --username "postgres" --drop
```

```
[postgres@ip-172-31-17-243 bin]$ psql -d pgconf -c "\d users"
Table "public.users"
 Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----+
 id   | integer |           | not null | nextval('pgosc_st_users_9bd652_id_seq'::regclass)
 name | character varying(100) |           | not null |
 email | character varying(255) |           | not null |
 price | double precision |           |          |
Indexes:
 "pgosc_st_users_9bd652_pkey" PRIMARY KEY, btree (id)

[postgres@ip-172-31-17-243 bin]$ pg-online-schema-change perform --alter-statement 'ALTER TABLE users ADD COLUMN "gender" BOOLEAN DEFAULT FALSE; ALTER TABLE users ADD COLUMN "last_name" varchar(20);' --dbname "pgconf" --host "localhost" --username "postgres" --drop > /tmp/multiple_alter.log 2>&1
[postgres@ip-172-31-17-243 bin]$ psql -d pgconf -c "\d users"
Table "public.users"
 Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----+
 id   | integer |           | not null | nextval( pgosc_st_users_27958a_id_seq ::regclass)
 name | character varying(100) |           | not null |
 email | character varying(255) |           | not null |
 price | double precision |           |          |
 gender | boolean |           |          | false
 last_name | character varying(20) |           |          |
Indexes:
 "pgosc_st_users_27958a_pkey" PRIMARY KEY, btree (id)
```

### *Renaming the column*

```
pg-online-schema-change perform --alter-statement 'ALTER TABLE
```

```
users RENAME COLUMN "price" TO "cost" --dbname "pgconf" --host  
"localhost" --username "postgres" --drop
```

```
[postgres@ip-172-31-17-243 bin]$ psql -d pgconf -c "vd users"  
Table "public.users"  
Column | Type | Collation | Nullable | Default  
-----+-----+-----+-----+-----  
id | integer | not null | nextval('pgosc_st_users_ceb1c0_id_seq'::regclass)  
name | character varying(100) | | |  
email | character varying(255) | | |  
cost | double precision | | |  
gender | boolean | | | false  
last_name | character varying(20) | | |  
Indexes:  
"pgosc_st_users_ceb1c0_pkey" PRIMARY KEY, btree (id)  
[postgres@ip-172-31-17-243 bin]$ pg-online-schema-change perform --alter-statement 'ALTER TABLE users RENAME COLUMN "cost" TO "price"' --dbname "pgconf" --host "localhost" --username "postgres" --drop >> /tmp/rename.log 2>l  
[postgres@ip-172-31-17-243 bin]$ psql -d pgconf -c "vd users"  
Table "public.users"  
Column | Type | Collation | Nullable | Default  
-----+-----+-----+-----+-----  
id | integer | not null | nextval('pgosc_st_users_ce7d14_id_seq'::regclass)  
name | character varying(100) | | |  
email | character varying(255) | | |  
price | double precision | | |  
gender | boolean | | | false  
last_name | character varying(20) | | |  
Indexes:  
"pgosc_st_users_ce7d14_pkey" PRIMARY KEY, btree (id)  
[postgres@ip-172-31-17-243 bin]$
```

## Difference between all above listed tools:

This is each tool's high-level features, use cases, and considerations to improve an application's operational performance and availability. We don't recommend one tool over another; you should choose the best tool depending on your use case.

## Monitoring

Online schema changes are very critical and have high impact on the live system, it's very important to have a robust monitoring system. Below are few areas that should be taken care during schema changes

Feature/Tool	pg-osc	pgroll	Flyway	pg_migrate	Liquibase	Blue-Green
<b>Primary Purpose</b>	Online schema changes	Zero-downtime schema changes	Database migration version control	Database migration	Database migration version control	Deployment strategy
<b>Database Support</b>	PostgreSQL only	PostgreSQL only	Multiple DBs	PostgreSQL only	Multiple DBs	Any
<b>Language/Platform</b>	Python	Go	Java	Node.js	Java	N/A
<b>Configuration Format</b>	SQL	YAML	SQL	JavaScript/SQL	XML, YAML, JSON, SQL	N/A
<b>Learning Curve</b>	Moderate	Moderate	Low	Low	High	Moderate
<b>Best For</b>	Large table alterations	Modern PostgreSQL apps	Java projects	Node.js projects	Enterprise apps	Any application
<b>Version Control</b>	No	Yes	Yes	Yes	Yes	N/A
<b>Rollback Support</b>	Yes	Yes	Limited	Yes	Yes	Yes
<b>Zero-Downtime</b>	Yes	Yes	No	No	No	Yes
<b>Enterprise Features</b>	Limited	Yes	Basic	Limited	Extensive	N/A
<b>Community Support</b>	Moderate	Growing	Large	Moderate	Large	N/A
<b>Release Year</b>	2015	2023	2010	2014	2006	N/A
<b>Scripting Support</b>	Limited	Yes	Yes	Yes	Yes	N/A
<b>Documentation Quality</b>	Good	Good	Excellent	Good	Excellent	N/A
<b>CI/CD Integration</b>	Limited	Good	Excellent	Good	Excellent	Excellent
<b>Cost</b>	Free	Free	Free/Commercial	Free	Free/Commercial	N/A

## Locks:

Since, locks are the most critical element in schema changes, ensure you've the most robust lock monitoring system/tools in place.

- "Active Session History in PostgreSQL: blocker and wait chain" by [Bertrand Drouvot](#) – this post describes the recursive CTE query `pg_ash_wait_chain.sql` that is useful for those who use the [pgsentinel](#) extension. The query is inspired by [Tanel Poder](#)'s script for Oracle.
- `locktree.sql` – query to display a tree of blocking sessions based on the information from `pg_locks` and `pg_stat_activity`, by Victor Yegorov.

## Query

```
with recursive activity as (
    select
        pg_blocking_pids(pid) blocked_by,
        *,
        age(clock_timestamp(), xact_start)::interval(0) as tx_age,
        -- "pg_locks.waitstart" - PG14+ only; for older
        versions: age(clock_timestamp(), state_change) as wait_age
        age(clock_timestamp(), (select max(l.waitstart) from
            pg_locks l where a.pid = l.pid))::interval(0) as wait_age
    from pg_stat_activity a
    where state is distinct from 'idle'
), blockers as (
    select
        array_agg(distinct c order by c) as pids
    from (
        select unnest(blocked_by)
        from activity
    ) as dt(c)
), tree as (
    select
        activity.*,
        1 as level,
        activity.pid as top_blocker_pid,
        array[activity.pid] as path,
        array[activity.pid]::int[] as all_blockers_above
    from activity, blockers
    where
        array[pid] <@ blockers.pids
        and blocked_by = '{}'::int[]
    union all
    select
        activity.*,
        tree.level + 1 as level,
        tree.top_blocker_pid,
        path || array[activity.pid] as path,
        tree.all_blockers_above || array_agg(activity.pid) over ()
    as all_blockers_above
    from activity, tree
    where
        not array[activity.pid] <@ tree.all_blockers_above
        and activity.blocked_by <> '{}'::int[]
        and activity.blocked_by <@ tree.all_blockers_above
)
select
```

```

pid,
blocked_by,
case when wait_event_type <> 'Lock' then replace(state, 'idle
in transaction', 'idletx') else 'waiting' end as state,
wait_event_type || ':' || wait_event as wait,
wait_age,
tx_age,
to_char(age(backend_xid), 'FM999,999,999,990') as xid_age,
to_char(2147483647 - age(backend_xmin), 'FM999,999,999,990')
as xmin_ttf,
datname,
username,
(select count(distinct t1.pid) from tree t1 where
array[tree.pid] <@ t1.path and t1.pid <> tree.pid) as blkd,
format(
'%s %s%s',
lpad('[' || pid::text || ']', 9, ' '),
repeat('.', level - 1) || case when level > 1 then ' ' end,
left(query, 1000)
) as query
from tree
order by top_blocker_pid, level, pid

```

\watch 10

## CPU/Memory:

For schema changes scenarios involving huge tables might be resource intensive operation especially with high CPU/Memory consumption. These kinds of operations might choke the system and be detrimental to the overall health of the system. It's highly recommended to have a robust monitoring system to monitor the CPU/Memory usage, the common metrics to observe are CPU Utilization, Load Average amongst the top OS processes.

## Disk Space:

In few of the schema changes operations/task, there would be a need to create a duplicate table, hence it's important to keep a close eye on the available disk space using your existing monitoring setup/tool.

## Live Demo

### Customer Case Studies

#### Adding check constraint on the existing column

#### Orders Table Structure

```
load=# \dt
      List of relations
 Schema |   Name    | Type | Owner
-----+-----+-----+
 public | customers | table | postgres
 public | order_items | table | postgres
 public | orders | table | postgres
(3 rows)

load=# \d orders
          Table "public.orders"
 Column |          Type          | Collation | Nullable | Default
-----+-----+-----+-----+-----+
 order_id | integer | not null | nextval('orders_order_id_seq'::regclass)
 customer_id | integer |          |          | CURRENT_TIMESTAMP
 order_date | timestamp without time zone |          |          |
 total_amount | numeric(10,2) |          |          |
 payment_method | character varying(20) |          |          |
 status | character varying(20) |          |          |
 shipping_address | text |          |          |
 delivery_date | date |          |          | CURRENT_TIMESTAMP
 created_at | timestamp without time zone |          |          |
 updated_at | timestamp without time zone |          |          | CURRENT_TIMESTAMP
Indexes:
 "orders_pkey" PRIMARY KEY, btree (order_id)
Foreign-key constraints:
 "orders_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customers(customer_id)
Referenced by:
 TABLE "order_items" CONSTRAINT "order_items_order_id_fkey" FOREIGN KEY (order_id) REFERENCES orders(order_id)
load=#
```

#### Orders data

```
load=# 
          order_id | customer_id | order_date | total_amount | payment_method | status | shipping_address | delivery_date | created_at | updated_at
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
      1 | 1 | 2025-02-15 09:03:08.912413 | 299.99 | credit_card | pending | 123 Main St, City | 2025-02-16 | 2025-02-15 09:03:08.912413 | 2025-02-15 09:03:08.912413
      2 | 2 | 2025-02-15 09:03:08.912413 | 159.50 | paypal | processing | 456 Oak Ave, Town | 2025-02-17 | 2025-02-15 09:03:08.912413 | 2025-02-15 09:03:08.912413
      3 | 3 | 2025-02-15 09:03:08.912413 | 499.99 | credit_card | delivered | 789 Pine Rd, Village | 2025-02-14 | 2025-02-15 09:03:08.912413 | 2025-02-15 09:03:08.912413
      4 | 4 | 2025-02-15 09:03:08.912413 | 499.99 |          | delivered | 789 Pine Rd, Village | 2025-02-14 | 2025-02-15 09:03:08.912413 | 2025-02-15 09:03:08.912413
(4 rows)
```

Observe payment\_method for the 4 order\_id its null value is added it is difficult to understand which is payment method?

How to resolve this issue?

1. Add default value (if anything is not added then default should be cash )or
2. Add check constraint (apart from given values anything comes then reject it along with null).

Consider you have application which needs to validate the new corrected data and also you need validation to avoid unknown values. To do this during business hours we need to do UPDATE to existing NULL to default values i.e. Cash and ALTER to add check constraint. Once it is done application should be able to see new value.

## What is impact of doing it during peak hours

1. Tables may get lock/query performance.
  2. CPU impact.
  3. Validation of application (if values are migrated properly or not).

## Test-1 Without schema change tool

1. Customer, Orders tables were occupied in one transaction.
  2. Update executed on orders table to update NULL status to cash and in waiting state
  3. Alter table to add constraint
  4. In parallel generating load using pgbench (which was locked and also alter was locked)

## Blocked:

**Once transaction completed it released the lock, altered the table and pgbench is running.**

## Test-2 With Schema change tool

1. Json file(to update/add check constraint)
  2. Execution of migration
  3. pgbench running update (which is locked for some very less time)

```
[postgres@ip-172-31-17-243 ~]$ cat 101_add_table_check_constraint.json
{
  "name": "101_add_table_check_constraint",
  "operations": [
    {
      "create_constraint": {
        "type": "check",
        "table": "orders",
        "name": "chk_valid_payment_method",
        "columns": [
          "payment_method"
        ],
        "check": "payment_method IN ('credit_card', 'paypal', 'bank_transfer', 'cash') AND payment_method is not null",
        "up": {
          "payment_method": "SELECT CASE WHEN payment_method is NULL THEN 'cash' ELSE payment_method END"
        },
        "down": {
          "payment_method": "payment_method"
        }
      }
    }
  ]
}
```

## Orders table (new column added by pgroll):

```

load=# \d orders
                                         Table "public.orders"
   Column      |      Type      | Collation | Nullable | Default
---+-----+-----+-----+-----+-----+
order_id      | integer      |           | not null | nextval('orders_order_id_seq'::regclass)
customer_id   | integer      |           |           | CURRENT_TIMESTAMP
order_date    | timestamp without time zone |           |           |
total_amount  | numeric(10,2) |           |           |
payment_method| character varying(20) |           |           |
status        | character varying(20) |           |           |
shipping_address| text |           |           |
delivery_date | date |           |           |
created_at    | timestamp without time zone |           |           |
updated_at    | timestamp without time zone |           |           |
_pgroll_new_payment_method| character varying(20) |           |           |
Indexes:
"orders_pkey" PRIMARY KEY, btree (order_id)
Check constraints:
"chk_valid_payment_method" CHECK ((._pgroll_new_payment_method::text = ANY (ARRAY['credit_card'::character varying, 'paypal'::character varying, 'bank_transfer'::character varying, 'cash'::character varying]::text[])) AND ._pgroll_new_payment_method IS NOT NULL) NOT VALID
Foreign-key constraints:
"orders_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customers(customer_id)
Referenced by:
TABLE "order_items" CONSTRAINT "order_items_order_id_fkey" FOREIGN KEY (order_id) REFERENCES orders(order_id)
Triggers:
_pgroll_trigger_orders_pgroll_new_payment_method BEFORE INSERT OR UPDATE ON orders FOR EACH ROW EXECUTE FUNCTION _pgroll_trigger_orders_pgroll_new_payment_method()
od() _pgroll_trigger_orders_payment_method BEFORE INSERT OR UPDATE ON orders FOR EACH ROW EXECUTE FUNCTION _pgroll_trigger_orders_payment_method()

```

## New column added by pgroll

```

load=# select payment_method,_pgroll_new_payment_method from orders;
          payment_method | _pgroll_new_payment_method
---+-----+
paypal | paypal
credit_card | credit_card
            | cash
paypal | paypal
credit_card | credit_card
            | cash
credit_card | credit_card
credit_card | credit_card
(8 rows)

```

```

load=# \dn
          List of schemas
   Name   | Owner
---+-----+
pgroll | postgres
public | pg_database_owner
public_101_add_table_check_constraint | postgres
(3 rows)

load=# \dt+ public_101_add_table_check_constraint.orders
          View "public_101_add_table_check_constraint.orders"
   Column      |      Type      | Collation | Nullable | Default | Storage | Description
---+-----+-----+-----+-----+-----+-----+
order_date    | timestamp without time zone |           |           |           | plain   |
customer_id   | integer      |           | not null |           | plain   |
total_amount  | numeric(10,2) |           |           |           | main    |
delivery_date | date |           |           |           | plain   |
payment_method| character varying(20) |           |           |           | extended|
shipping_address| text |           |           |           | extended|
status        | character varying(20) |           |           |           | extended|
order_id      | integer      |           |           |           | plain   |
created_at    | timestamp without time zone |           |           |           | plain   |
updated_at    | timestamp without time zone |           |           |           | plain   |
View definition:
SELECT orders.order_date,
       orders.customer_id,
       orders.total_amount,
       orders.delivery_date,
       orders._pgroll_new_payment_method AS payment_method,
       orders.shipping_address,
       orders.status,
       orders.order_id,
       orders.created_at,
       orders.updated_at
  FROM orders;
Options: security_invoker=true

```

## After complete migration:

```
[postgres@ip-172-31-17-243 ~]$ pgroll complete --postgres-url postgres://postgres:postgres@localhost:5432/load?sslmode=disable 101_add_table_check_constraint.json
[postgres@ip-172-31-17-243 ~]$ psql -d load
psql (15.9)
Type "help" for help.

load=# \d orders
              Table "public.orders"
   Column   |      Type       | Collation | Nullable | Default
-----+-----+-----+-----+-----+
order_id | integer |          | not null | nextval('orders_order_id_seq'::regclass)
customer_id | integer |          |          | CURRENT_TIMESTAMP
order_date | timestamp without time zone |          |          | CURRENT_TIMESTAMP
total_amount | numeric(10,2) |          |          | CURRENT_TIMESTAMP
item_id | character varying(20) |          |          |
shipping_address | text |          |          |
delivery_date | date |          |          |
created_at | timestamp without time zone |          |          | CURRENT_TIMESTAMP
updated_at | timestamp without time zone |          |          | CURRENT_TIMESTAMP
payment_method | character varying(20) |          |          |
Indexes:
"orders_pkey" PRIMARY KEY, btree (order_id)
Check constraints:
"payment_method" CHECK ((payment_method::text = ANY (ARRAY['credit_card'::character varying, 'paypal'::character varying, 'bank_transfer'::character varying, 'cash'::character varying])) AND payment_method IS NOT NULL)
Foreign-key constraints:
"orders_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customers(customer_id)
Referenced by:
TABLE "order_items" CONSTRAINT "order_items_order_id_fkey" FOREIGN KEY (order_id) REFERENCES orders(order_id)
```

## Status of last migration

```
[postgres@ip-172-31-17-243 ~]$ pgroll status --postgres-url postgres://postgres:postgres@localhost:5432/load?sslmode=disable 101_add_table_check_constraint.json
{
  "schema": "public",
  "version": "101_add_table_check_constraint",
  "status": "Complete"
}
[postgres@ip-172-31-17-243 ~]$
```

## Best Practices

- **Splitting the work into batches** - Consider splitting the work into batches, each one being a separate transaction. If you're working in the OLTP context (mobile or web apps), the batch size should be determined appropriately
- **Take care or Tune VACUUMing** - Tune autovacuum and/or consider using explicit VACUUM calls after some number of batches processed.
- **Use extended lock analysis:** As explained in the monitoring section
- **Use CONCURRENTLY** when adding indexes, but that doesn't work inside transactions.
- When adding a constraint like NOT NULL, Postgres needs to first check that there are no NULLs in the table, and it has to do that while holding the lock. You can work around it by adding a **CHECK CONSTRAINT with NOT VALID**, which means that the constraint is applied to new rows but not to old ones. Then you can run VALIDATE later, which doesn't need the ACCESS EXCLUSIVE lock because it assumes new rows are respecting the constraint.
- **Tune checkpoint** - If not well tuned checkpoints may occur very often during such a massive operation. as an extra protection measure, tune the

checkpoint so that even if a massive change happens, our database's negative effect is not so acute ["Basics of Tuning Checkpoints](#)

- Use Lock-retry methodology at the SQL level - **lock\_timeout**, to acquire a lock on the DB objects that are subject to change and do not implement some kind of retry logic.
- **DML never should go after DDL** unless they both deal with some freshly created table, It is usually wise to split DDL and DML activities into separate transactions / migration steps
- Increase **Statement\_timeout** parameter value, if you have both Production and Non-Production environments using same values of statement\_timeout, it's quite possible that the smaller tables are, the faster queries are executed. This can easily lead to a situation when a timeout is reached only on production.

## Conclusion:

Online schema changes in PostgreSQL require careful planning and execution to minimize downtime and maintain data integrity. The most effective approach combines multi-phase deployments with backward-compatible changes, utilizing PostgreSQL's built-in features like transactional DDL and concurrent index creation. While some operations like adding nullable columns and creating indexes concurrently are relatively safe, others such as changing column types or adding NOT NULL constraints require special attention and potentially more complex migration strategies. Success depends on following best practices: implementing changes incrementally, maintaining backward compatibility, thorough testing with production-scale data, and having clear rollback plans. Tools like Liquibase and Flyway can help manage these changes, but understanding the underlying implications of each schema modification remains crucial for maintaining system reliability during migrations.

## Disclaimer:

**NO WARRANTY:** The provided information is for guidance only. Implementation is at your own risk. Neither the authors nor their organizations are liable for any damages or losses resulting from the use of these procedures.

**RECOMMENDATION:** Always test thoroughly in non-production environments first and engage qualified database administrators for critical production changes.

## **References:**

[Lock Types](#)

[Pgroll](#)

[Pg OSC](#)

[Common DB Schema changes mistakes](#)

[PostgreSQL Locks](#)